

AlexNet

Leon Friedrich, INFM



**Elektrotechnik, Medizintechnik
und Informatik**

Agenda

- Convolution
- AlexNet
 - Why AlexNet?
 - Architecture
- Practical Part
 - Visualization of learned Kernels
 - Visualization of Convolution

Convolution I

- Mathematical operation: $*$
- 2D convolution in image processing:
 - Image $I[u, v]$
 - Kernel K with size $[2a + 1, 2b + 1]$
 - Output $O[u, v]$

	-1	0	1
-1	1	2	3
0	4	5	6
1	7	8	9

$$O[u, v] = I[u, v] * K = \sum_{s=-a}^a \sum_{t=-b}^b I[u - s, v - t] K[s, t]$$

Convolution II

$$O[u, v] = I[u, v] * K = \sum_{s=-a}^a \sum_{t=-b}^b I[u-s, v-t] K[s, t]$$

1	2	3
4	5	6
7	8	9

K

flip

9	8	7
6	5	4
3	2	1

9	8	7	0	0	0	0
6	5	4	0	0	0	0
3	2	1	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$I[u, v]$

Convolution III

9	8	7	0	0	0	0
6	5	4	0	0	0	0
3	2	1	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$I[u, v]$

$$O[1,1] = 9 \cdot 0 + 8 \cdot 0 + 7 \cdot 0 + 6 \cdot 0 + 5 \cdot 0 + 4 \cdot 0 + 3 \cdot 0 + 2 \cdot 0 + 1 \cdot 0 = 0$$

	0					

$O[u, v]$

Convolution III

0	9	8	7	0	0	0
0	6	5	4	0	0	0
0	3	2	1	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$I[u, v]$

$$O[1,2] = 9 \cdot 0 + 8 \cdot 0 + 7 \cdot 0 + 6 \cdot 0 + 5 \cdot 0 + 4 \cdot 0 + 3 \cdot 0 + 2 \cdot 0 + 1 \cdot 0 = 0$$

	0	0				

$O[u, v]$

Convolution III

0	0	9	8	7	0	0
0	0	6	5	4	0	0
0	0	3	2	1	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$I[u, v]$

$$O[1,3] = 9 \cdot 0 + 8 \cdot 0 + 7 \cdot 0 + 6 \cdot 0 + 5 \cdot 0 + 4 \cdot 0 + 3 \cdot 0 + 2 \cdot 0 + 1 \cdot 0 = 0$$

	0	0	0			

$O[u, v]$

Convolution III

0	0	0	9	8	7	0
0	0	0	6	5	4	0
0	0	0	3	2	1	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$I[u, v]$

$$O[1,4] = 9 \cdot 0 + 8 \cdot 0 + 7 \cdot 0 + 6 \cdot 0 + 5 \cdot 0 + 4 \cdot 0 + 3 \cdot 0 + 2 \cdot 0 + 1 \cdot 0 = 0$$

	0	0	0	0		

$O[u, v]$

Convolution III

0	0	0	0	9	8	7
0	0	0	0	6	5	4
0	0	0	0	3	2	1
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$I[u, v]$

$$O[1,5] = 9 \cdot 0 + 8 \cdot 0 + 7 \cdot 0 + 6 \cdot 0 + 5 \cdot 0 + 4 \cdot 0 + 3 \cdot 0 + 2 \cdot 0 + 1 \cdot 0 = 0$$

	0	0	0	0	0	

$O[u, v]$

Convolution III

0	0	0	0	0	0	0
9	8	7	0	0	0	0
6	5	4	0	0	0	0
3	2	1	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$I[u, v]$

$$O[2,1] = 9 \cdot 0 + 8 \cdot 0 + 7 \cdot 0 + 6 \cdot 0 + 5 \cdot 0 + 4 \cdot 0 + 3 \cdot 0 + 2 \cdot 0 + 1 \cdot 1 = 1$$

	0	0	0	0	0	
	0					

$O[u, v]$

Convolution III

0	0	0	0	0	0	0
0	9	8	7	0	0	0
0	6	5	4	0	0	0
0	3	2	1	1	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0


$I[u, v]$

$$O[2,2] = 9 \cdot 0 + 8 \cdot 0 + 7 \cdot 0 + 6 \cdot 0 + 5 \cdot 0 + 4 \cdot 0 + 3 \cdot 0 + 2 \cdot 0 + 1 \cdot 1 = 1$$

	0	0	0	0	0	
	0	1				

$O[u, v]$

Convolution IV



Convolution V

-1	-2	-1
0	0	0
1	2	1

*




K

$I[u, v]$

$O[u, v]$

AlexNet - Why AlexNet?


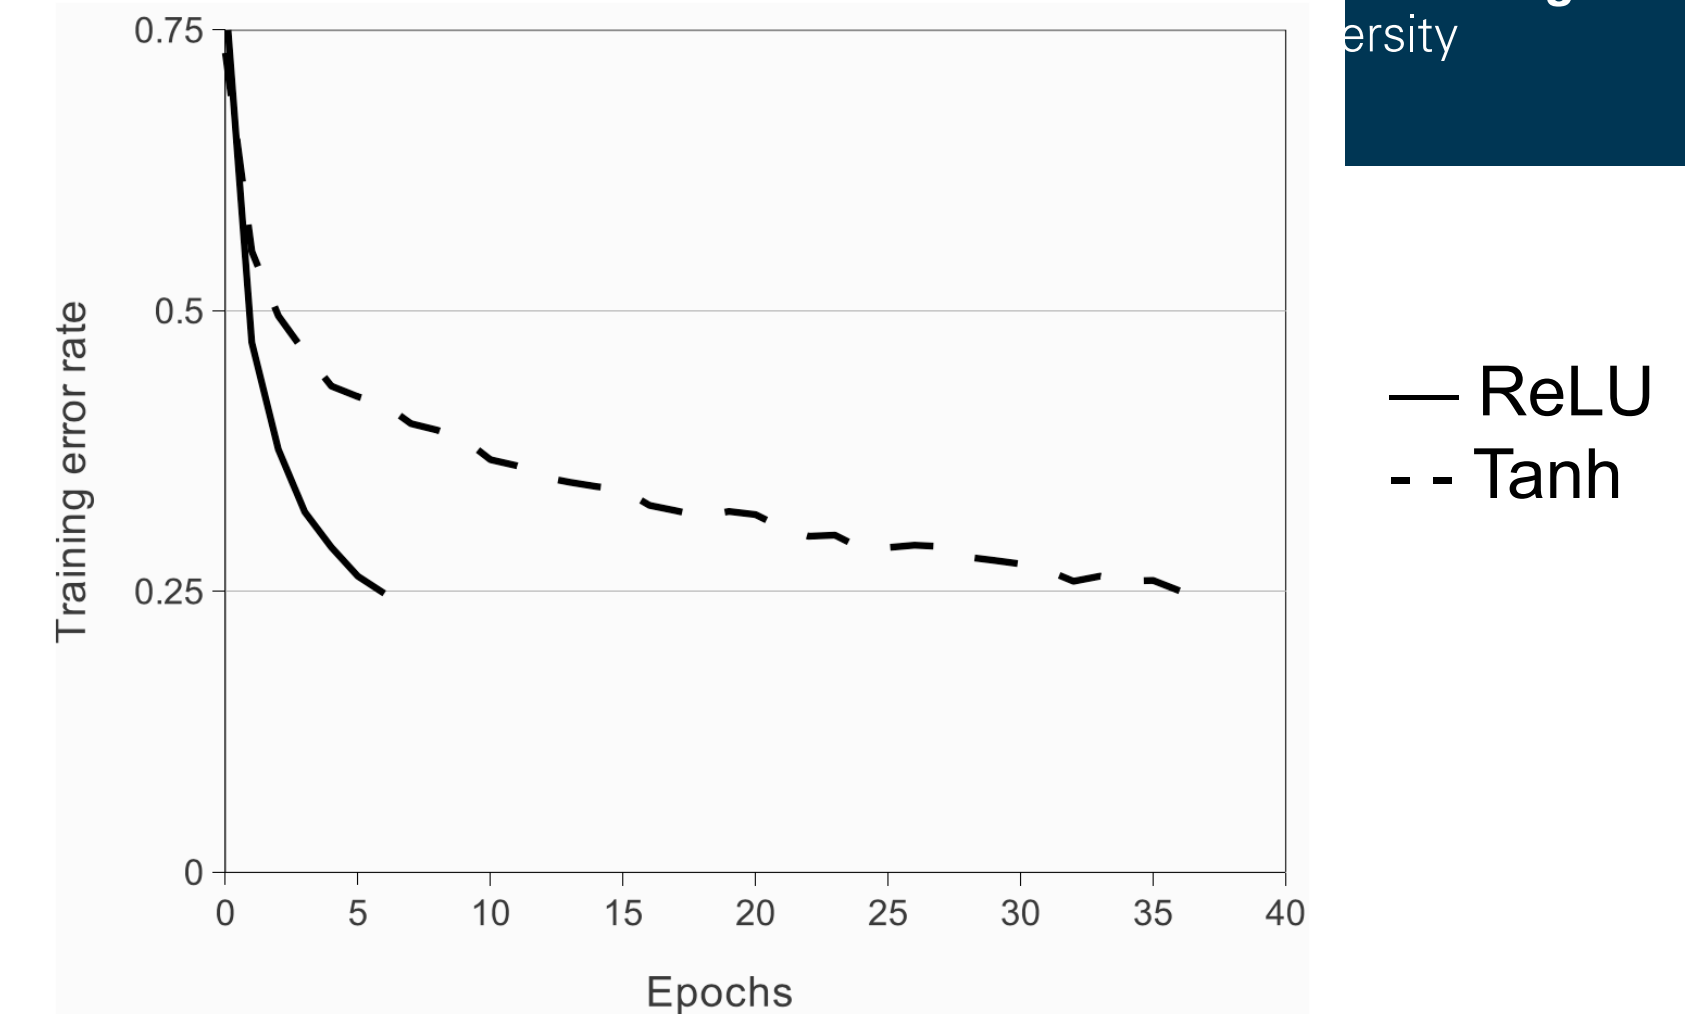
- Alex Krizhevsky et al. (2012)
- Winner ILSVRC-2012
 - Top-5 test error rate: 15.3%
- Key to success: Faster training



By Gkrusze - Own work, CC BY-SA 4.0,
<https://commons.wikimedia.org/w/index.php?curid=69750373>

AlexNet - Architecture I


- 5 convolutional, 3 fully-connected
- Multiple GPUs
- ReLU: $f(x) = \max(0, x)$




AlexNet - Reduce Overfitting

- 60 Mio. parameters
- Data Augmentation
- Dropout
- Overlapping Pooling

Non-overlapping pooling



Overlapping pooling



Source: <https://bskyvision.com/421>

Practical Part - AlexNet

Manually implemented:



```
AlexNet(  
  (features): Sequential(  
    (0): Conv2d(3, 96, kernel_size=(11, 11), stride=(4, 4))  
    (1): ReLU()  
    (2): LocalResponseNorm(5, alpha=0.0001, beta=0.75, k=2)  
    (3): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (4): Conv2d(96, 256, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))  
    (5): ReLU()  
    (6): LocalResponseNorm(5, alpha=0.0001, beta=0.75, k=2)  
    (7): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (8): Conv2d(256, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): ReLU()  
    (10): Conv2d(384, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (11): ReLU()  
    (12): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (13): ReLU()  
    (14): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (classifier): Sequential(  
    (0): Dropout(p=0.5, inplace=True)  
    (1): Linear(in_features=256, out_features=4096, bias=True)  
    (2): ReLU()  
    (3): Dropout(p=0.5, inplace=True)  
    (4): Linear(in_features=4096, out_features=4096, bias=True)  
    (5): ReLU()  
    (6): Linear(in_features=4096, out_features=8, bias=True)  
  )  
)
```

Pytorch version:



```
AlexNet(  
  (features): Sequential(  
    (0): Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2))  
    (1): ReLU(inplace=True)  
    (2): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (3): Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2))  
    (4): ReLU(inplace=True)  
    (5): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
    (6): Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (7): ReLU(inplace=True)  
    (8): Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (9): ReLU(inplace=True)  
    (10): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
    (11): ReLU(inplace=True)  
    (12): MaxPool2d(kernel_size=3, stride=2, padding=0, dilation=1, ceil_mode=False)  
  )  
  (avgpool): AdaptiveAvgPool2d(output_size=(6, 6))  
  (classifier): Sequential(  
    (0): Dropout(p=0.5, inplace=False)  
    (1): Linear(in_features=9216, out_features=4096, bias=True)  
    (2): ReLU(inplace=True)  
    (3): Dropout(p=0.5, inplace=False)  
    (4): Linear(in_features=4096, out_features=4096, bias=True)  
    (5): ReLU(inplace=True)  
    (6): Linear(in_features=4096, out_features=1000, bias=True)  
  )  
)
```

Accuracy and Loss



Manual Implementation



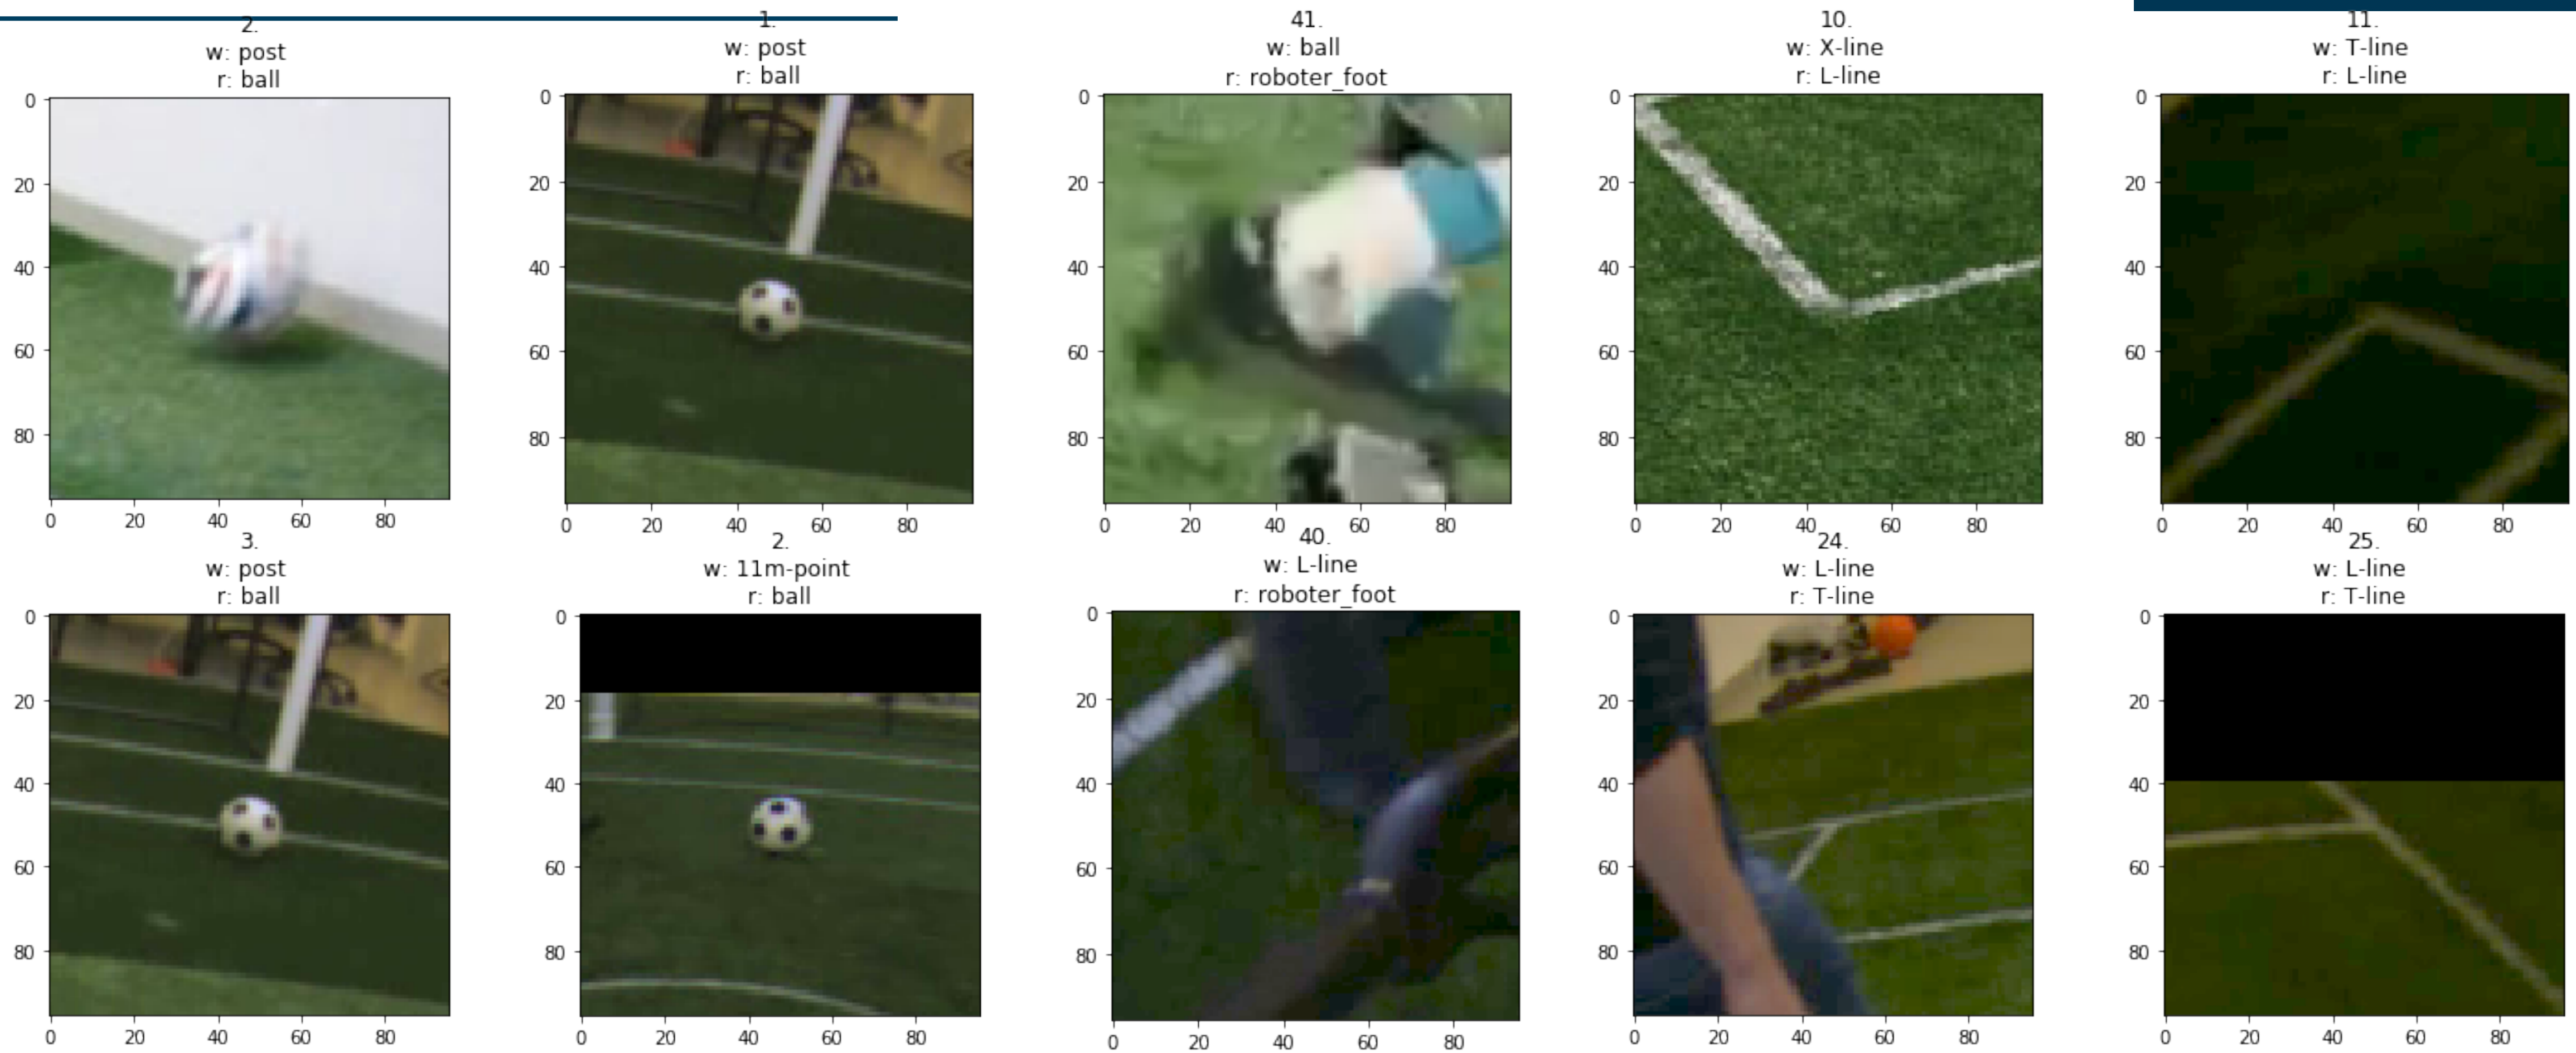
Pytorch not pretrained



Pytorch pretrained



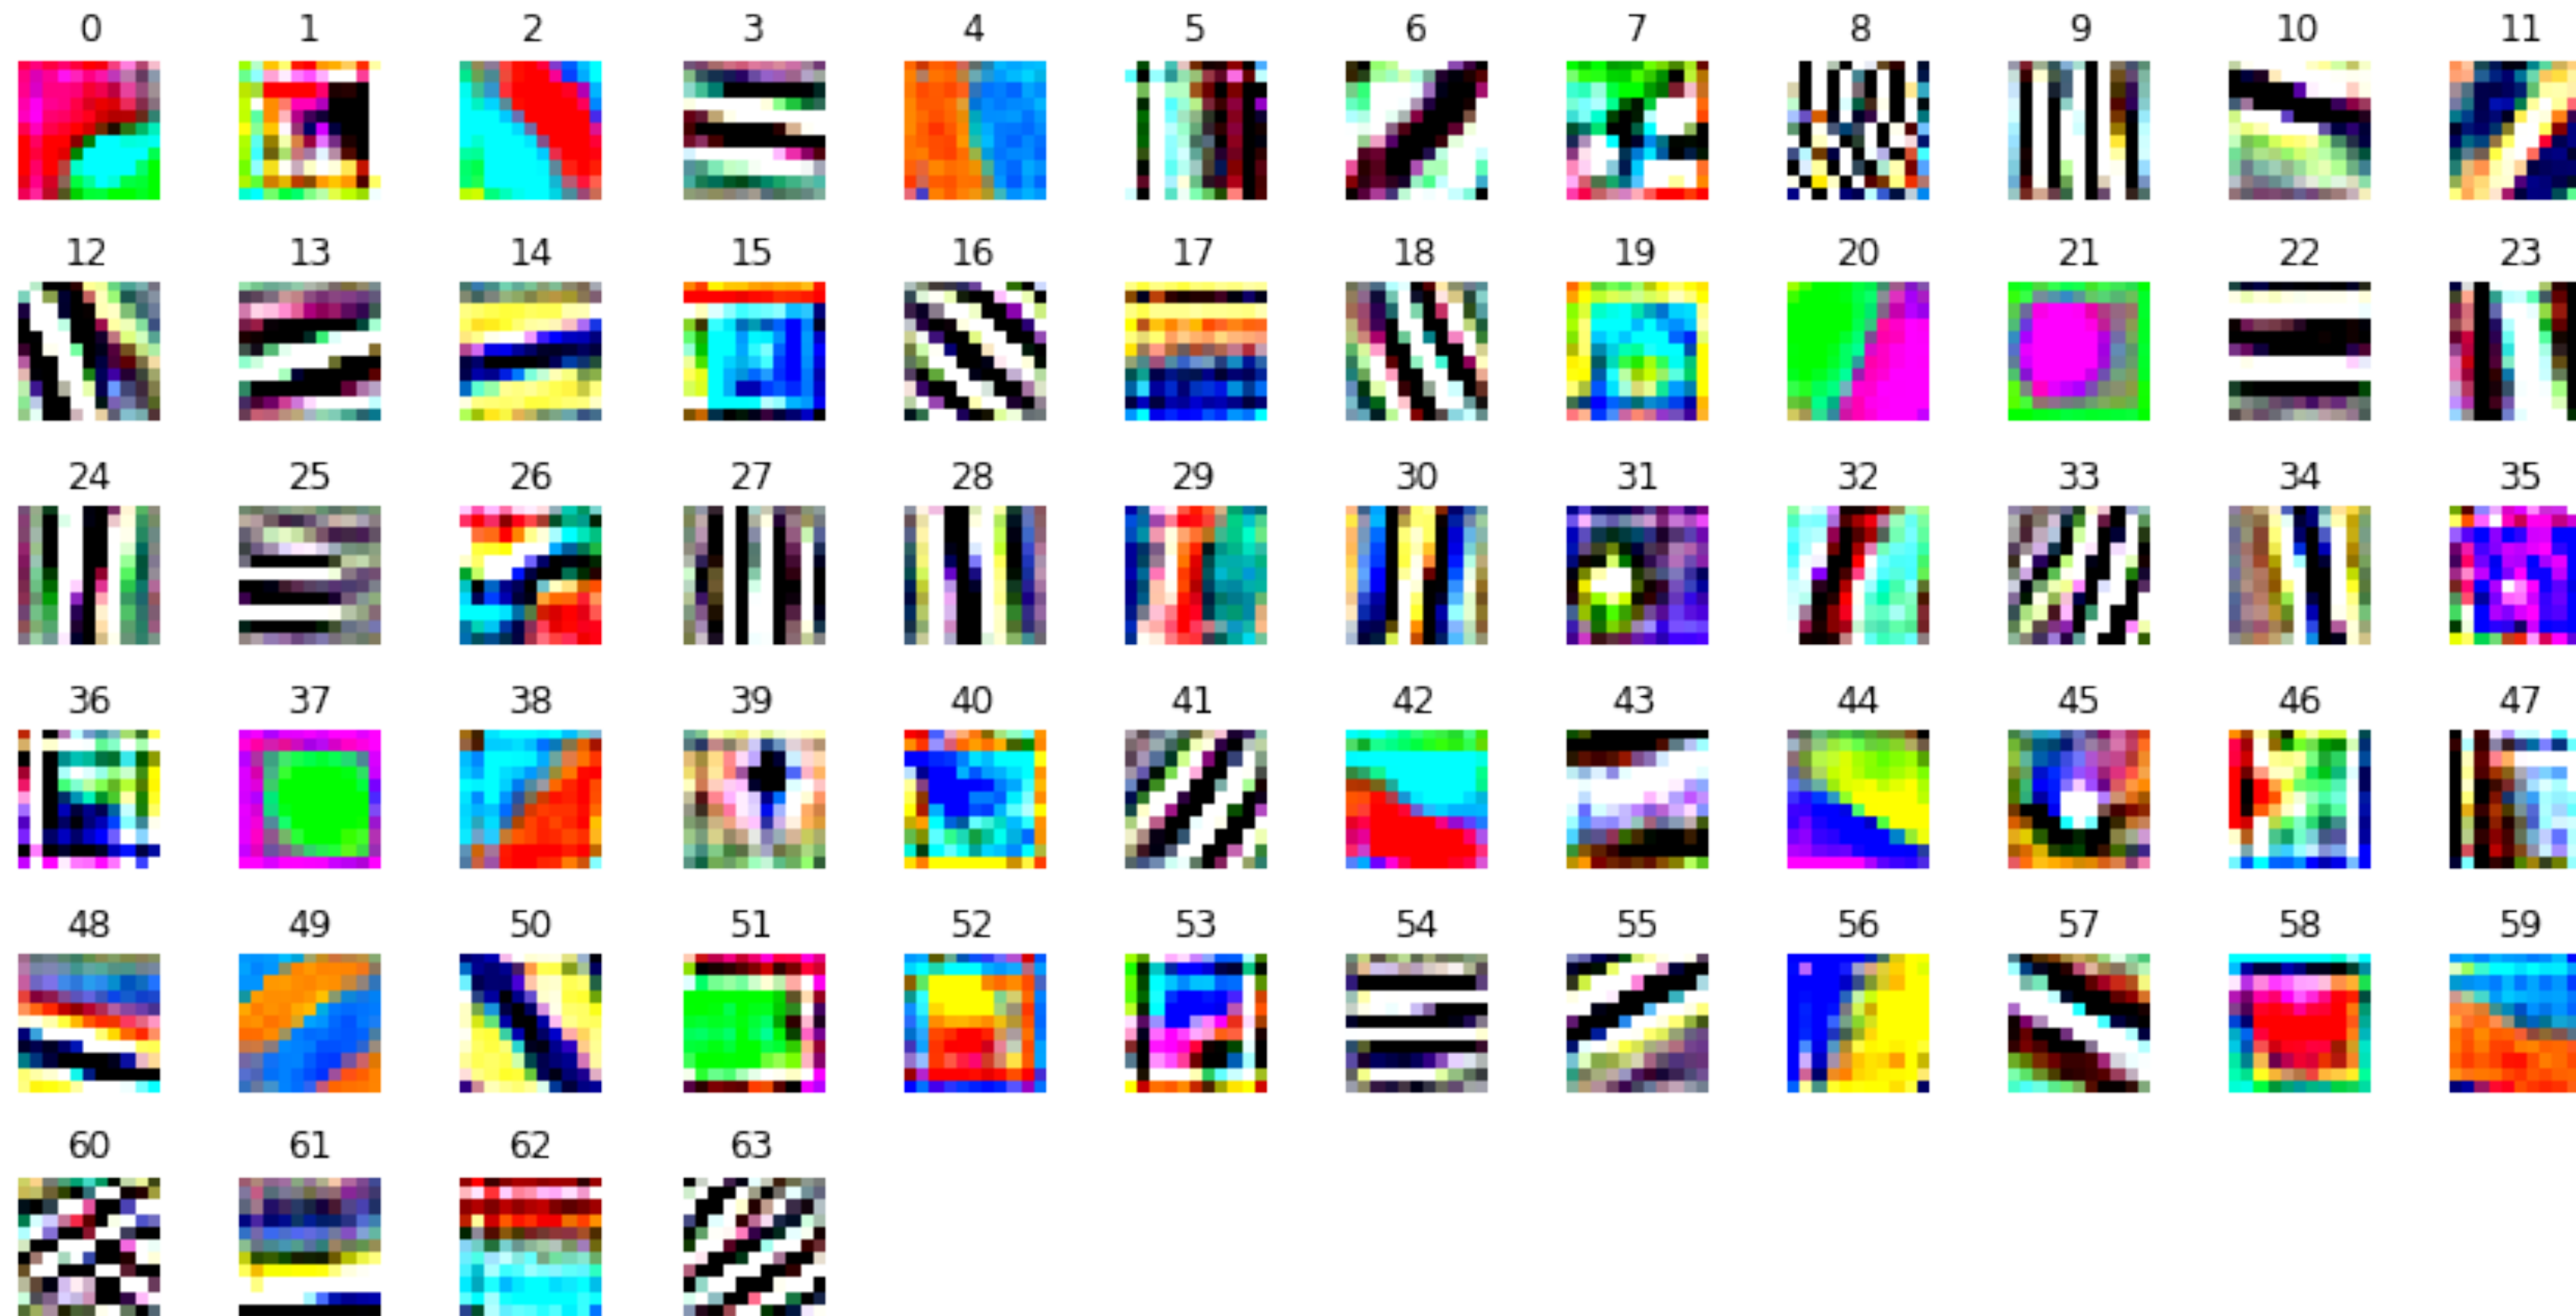
Wrongly Classified Images



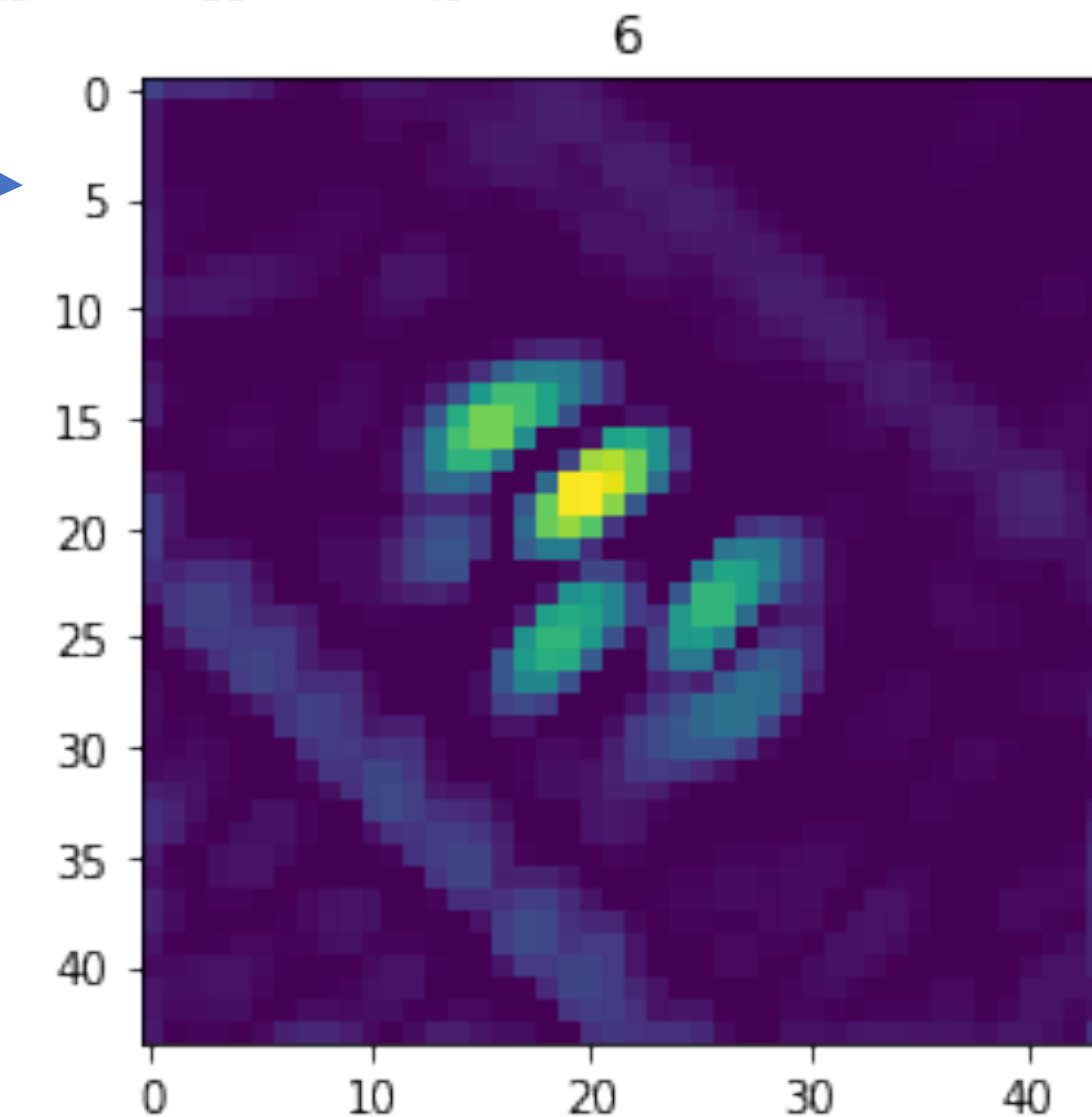


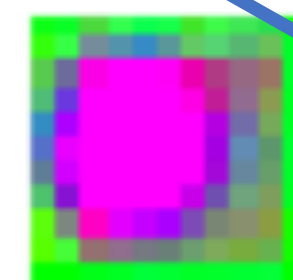

Learned Kernels - Not Pretrained



Learned Kernels - Pretrained



Visualization of Convolution



Thank you!

Sources

Paper:

<https://papers.nips.cc/paper/4824-imagenet-classification-with-deep-convolutional-neural-networks.pdf>

Gitlab Repo with slides, notebook and paper:

<https://gitlab.iz.hs-offenburg.de/lfriedri/deep-learning-alexnet.git>

Further sources:

<https://en.wikipedia.org/wiki/AlexNet>

<https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

https://machinelearningguru.com/computer_vision/basics/convolution/image_convolution_1.html

[https://en.wikipedia.org/wiki/Kernel_\(image_processing\)](https://en.wikipedia.org/wiki/Kernel_(image_processing))